

6.1 Training Neural Networks

Machine Learning 1: Foundations

Marius Kloft (TUK)

Recap

Artificial neural networks (ANN)


- ▶ Key advantage over SVM, logistic regression, and friends: can **learn a good representation** of the data,

$$\min_{b \in \mathbb{R}, \mathbf{w} \in \mathbb{R}^d, \phi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \ln \left(1 + \exp(-y_i(\mathbf{w}^\top \phi(\mathbf{x}_i) + b)) \right).$$

Need to restrict search space of ϕ !

Idea: design ϕ similar to our brain

- ▶ multiple neurons in multiple layers with feed-forward connections
- ▶ $\phi_W(\mathbf{x}_i) := \sigma(W_L^\top \dots \sigma(W_1^\top \cdot \mathbf{x}_i) \dots)$
- ▶ optimize over $W = (W_1, \dots, W_L)$!



How to train ANNs?

Contents of this Class

1 Training Neural Networks

2 Deep Learning

1 Training Neural Networks

2 Deep Learning

How to Train (Deep) ANNs?

In the same way as we trained the SVM:
(stochastic) gradient descent!

Recall the ANN optimization problem:

$$\min_{\mathbf{w}, W} \underbrace{\frac{1}{2} \|\mathbf{w}\|^2 + \frac{1}{2} \sum_{l=1}^L \|W_l\|_{\text{Fro}}^2 + C \sum_{i=1}^n \ln \left(1 + \exp \left(-y_i \mathbf{w}^\top \phi_W(\mathbf{x}_i) \right) \right)}_{=: F(\mathbf{w}, W)}$$

How to compute the gradient of F ?

For the sake of simplicity, we focus on discussing how to train *fully connected ANNs* (not CNNs).

The Gradient of F With Respect to \mathbf{w} is Simple:

The function

$$g(x) = \ln(1 + \exp(x))$$

has the derivative

$$g'(x) = \frac{\exp(x)}{1 + \exp(x)} = \frac{1}{1 + \exp(-x)}.$$

Thus, by the chain rule:

$$\begin{aligned}\nabla_{\mathbf{w}} F(\mathbf{w}, W) &= \mathbf{w} + C \sum_{i=1}^n \nabla_{\mathbf{w}} \ln \left(1 + \exp \left(-y_i \mathbf{w}^\top \phi_W(\mathbf{x}_i) \right) \right) \\ &= \mathbf{w} - C \sum_{i=1}^n \frac{y_i \phi_W(\mathbf{x}_i)}{1 + \exp(y_i \mathbf{w}^\top \phi_W(\mathbf{x}_i))}\end{aligned}$$

But how to compute the gradient of F with respect to W ?

Gradient of F With Respect to $W = (W_1, \dots, W_L)$

Analogously, we have, for all $l = 1, \dots, L$:

$$\begin{aligned}\nabla_{W_l} F(\mathbf{w}, W) &= W_l + C \sum_{i=1}^n \nabla_{W_l} \ln \left(1 + \exp \left(- y_i \mathbf{w}^\top \phi_W(\mathbf{x}_i) \right) \right) \\ &= W_l - C \sum_{i=1}^n \frac{y_i \mathbf{w}^\top \nabla_{W_l} \phi_W(\mathbf{x}_i)}{1 + \exp(y_i \mathbf{w}^\top \phi_W(\mathbf{x}_i))},\end{aligned}$$

where we applied the chain rule.

From now on, denote the ij th entry of W_l by w_{ijl} .

Given a data point \mathbf{x} , how to compute $\nabla_{w_{ijl}} \phi_W(\mathbf{x})$?

Computing $\nabla_{w_{ijl}} \phi_W(\mathbf{x})$

We have:

$$\nabla_{w_{ijl}} \phi_W(\mathbf{x}) = \nabla_{w_{ijl}} \sigma \left(\underbrace{W_L^\top \sigma \left(\underbrace{\dots \sigma \left(\underbrace{W_1^\top \mathbf{v}_0}_{=\mathbf{u}_1} \right) \dots}_{=\mathbf{v}_1} \right)}_{\mathbf{u}_L} \right)_{\mathbf{v}_L}$$

Need to compute a gradient of a **nested** function!

Idea: Chain rule

$$\nabla_{w_{ijl}} \phi_W(\mathbf{x}) = \frac{\partial \mathbf{v}_L}{\partial w_{ijl}} = \frac{\partial \mathbf{v}_L}{\partial \mathbf{u}_L} \cdot \frac{\partial \mathbf{u}_L}{\partial \mathbf{v}_{L-1}} \cdot \frac{\partial \mathbf{v}_{L-1}}{\partial \mathbf{u}_{L-1}} \dots \frac{\partial \mathbf{u}_{l+1}}{\partial \mathbf{v}_l} \cdot \frac{\partial \mathbf{v}_l}{\partial \mathbf{u}_l} \cdot \frac{\partial \mathbf{u}_l}{\partial w_{ijl}}$$

①
②
①
②
①
③

Three Terms Occur by the Chain Rule:

For all $l = 1, \dots, L$:

① $\frac{\partial \mathbf{v}_l}{\partial \mathbf{u}_l}$

② $\frac{\partial \mathbf{u}_l}{\partial \mathbf{v}_{l-1}}$

③ $\frac{\partial \mathbf{u}_l}{\partial w_{ijl}}$

We need to compute all of them!

First Term

We compute the first term as:

$$\textcircled{1} \quad \frac{\partial \mathbf{v}_l}{\partial \mathbf{u}_l} = \frac{\partial \sigma(\mathbf{u}_l)}{\partial \mathbf{u}_l} = \frac{\partial \max(0, \mathbf{u}_l)}{\partial \mathbf{u}_l} = \Theta(\mathbf{u}_l),$$

where

$$\begin{array}{ccc} \mathbb{R} & \rightarrow & \mathbb{R} \\ \Theta : x & \mapsto & \begin{cases} 0 & \text{if } x \leq 0 \\ 1 & \text{otherwise} \end{cases} \end{array}$$

is the **heavyside function**, which, for a vector $\mathbf{x} = (x_1, \dots, x_d)^\top \in \mathbb{R}^d$, is defined elementwise:

$$\Theta(\mathbf{x}) := \begin{pmatrix} \Theta(x_1) \\ \vdots \\ \Theta(x_d) \end{pmatrix}.$$

Second Term

We compute the second term as:

$$\textcircled{2} \quad \frac{\partial \mathbf{u}_l}{\partial \mathbf{v}_{l-1}} = \frac{\partial (W_l^\top \mathbf{v}_{l-1})}{\partial \mathbf{v}_{l-1}} = W_l^\top$$

Third Term

We compute the third term as:

$$\textcircled{3} \quad \frac{\partial \mathbf{u}_l}{\partial \mathbf{w}_{ijl}} = \frac{\partial (\mathbf{W}_l^\top \mathbf{v}_{l-1})}{\partial \mathbf{w}_{ijl}} = \left(\frac{\partial \sum_k \mathbf{w}_{kk'l} v_{k,l-1}}{\partial \mathbf{w}_{ijl}} \right)_{k'} = v_{i,l-1} \mathbf{e}_j$$

where

- ▶ $v_{i,l-1}$ denotes the i th entry of \mathbf{v}_{l-1}
- ▶ \mathbf{e}_j is a unit vector with entries zero everywhere except in the j th component.

Putting Things Together

Our chain rule formula from Slide 7 thus translates into:

$$\begin{aligned}\nabla_{w_{ijl}} \phi_W(\mathbf{x}) &= \frac{\partial \mathbf{v}_L}{\partial \mathbf{u}_L} \cdot \frac{\partial \mathbf{u}_L}{\partial \mathbf{v}_{L-1}} \cdot \frac{\partial \mathbf{v}_{L-1}}{\partial \mathbf{u}_{L-1}} \cdots \frac{\partial \mathbf{u}_{l+1}}{\partial \mathbf{v}_l} \cdot \frac{\partial \mathbf{v}_l}{\partial \mathbf{u}_l} \cdot \frac{\partial \mathbf{u}_l}{\partial w_{ijl}} \\ &= \Theta(\mathbf{u}_L) W_L^\top \Theta(\mathbf{u}_{L-1}) \cdots W_{l+1}^\top \Theta(\mathbf{u}_l) \mathbf{v}_{i,l-1} \mathbf{e}_j\end{aligned}$$

How to code up the computation of

$$\nabla_{w_{ijl}} \phi_W(\mathbf{x}) \quad \forall i, j, l$$

in an efficient algorithm?

Backpropagation Algorithm

Given an input \mathbf{x} , we first compute all variables \mathbf{u}_l and \mathbf{v}_l :

Forward propagation

- 1: initialize $\mathbf{v}_0 := \mathbf{x}$
- 2: **for** $l = 1 : (L - 1)$ **do**
- 3: $\mathbf{u}_l := \mathbf{W}_l^\top \mathbf{v}_{l-1}$
- 4: $\mathbf{v}_l := \sigma(\mathbf{u}_l)$
- 5: **end for**

Then, we compute the gradient via the chain rule:

Backward propagation

- 1: initialize $\delta_L := \Theta(\mathbf{u}_L)$
- 2: $\nabla_{\mathbf{w}_{jl}} \phi_W(\mathbf{x}) := \delta_L v_{i,L-1} \mathbf{e}_j \quad \forall i, j$
- 3: **for** $l = (L - 1) : 1$ **do**
- 4: $\delta_l := \delta_{l+1} \mathbf{W}_{l+1}^\top \Theta(\mathbf{u}_l)$
- 5: $\nabla_{\mathbf{w}_{jl}} \phi_W(\mathbf{x}) := \delta_l v_{i,l-1} \mathbf{e}_j \quad \forall i, j$
- 6: **end for**

Conclusion

How to train ANNs?

- ▶ Stochastic gradient descent

How to compute gradient?

- ▶ ANN is a nested function
- ▶ Thus we compute the gradient via the chain rule
- ▶ Lead to a recursive algorithm: backpropagation

Outlook

Advanced training algorithms:

- ▶ Adagrad
- ▶ Adam
- ▶ Nesterov momentum