

7.2 Loss View

Machine Learning 1: Foundations

Marius Kloft (TUK)

- 1 The Problem: Overfitting
- 2 Unifying View
- 3 The Solution: Regularization
- 4 Regularization for Deep Learning

Recall:

SVM:

$$\min_{b, \mathbf{w}} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \max \left(0, 1 - y_i (\langle \mathbf{w}, \phi_k(\mathbf{x}_i) \rangle + b) \right)$$

↑

kernel feature map of a kernel k

Logistic Regression (LR):

$$\min_{b, \mathbf{w}} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \ln \left(1 + \exp(-y_i (\langle \mathbf{w}, \phi_k(\mathbf{x}_i) \rangle + b)) \right)$$

ANN:

$$\min_{b, \mathbf{w}, \mathbf{W}} \frac{1}{2} \|\mathbf{w}\|^2 + \frac{1}{2} \sum_{l=1}^L \|\mathbf{W}_l\|_{\text{Fro}}^2 + C \sum_{i=1}^n \ln(1 + \exp(-y_i (\mathbf{w}^\top \phi_W(\mathbf{x}_i) + b)))$$

All these methods can be unified into a single equation.

For simplicity, we had introduced ANNs without bias b in the ANN class. Here we use a bias b , which makes a lot sense, for the same reasons as it makes sense also in SVMs and LR.

Unifying View

Unifying formulation of linear, kernel, and neural learning

$$\min_{[W,] b, \mathbf{w}} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \ell(y_i(\langle \mathbf{w}, \phi(\mathbf{x}_i) \rangle + b)) \left[+ \frac{1}{2} \sum_{l=1}^L \|W_l\|_{\text{Fro}}^2 \right],$$

where

- ▶ $\ell(t) := \max(0, 1 - t)$ for SVM (“hinge loss”)
- ▶ $\ell(t) := \ln(1 + \exp(-t))$ for LR and ANN (“logistic loss”)

and

- ▶ $\phi := \text{id}$ for linear SVM and linear LR
- ▶ $\phi := \phi_k$ for kernel SVM and kernel LR
- ▶ $\phi := \phi_W$ for ANN.

The terms in brackets apply only for ANNs.

For a test point \mathbf{x} , we predict $f(\mathbf{x}) := \text{sign}(\langle \mathbf{w}, \phi(\mathbf{x}) \rangle + b)$.

Five Popular Learning Machines in One Equation

The following table summarizes the result of the previous slide:

<div>map ϕ</div> <div>loss ℓ</div>	id	ϕ_k	ϕ_w
hinge	linear SVM	kernel SVM	— ¹
logistic	linear LR	kernel LR	ANN

¹ The hinge loss is theoretically possible but uncommon in neural networks.

The Loss


The unifying equation contains,
for every training example (\mathbf{x}_i, y_i) ,
a term

$$\underbrace{\ell(y_i \langle \mathbf{w}, \phi(\mathbf{x}) \rangle + b)}_{=: t_i},$$

the “loss” of the i th example.

Interpretation:


- ▶ In the unifying formulation, we minimize the loss $\sum_i \ell(t_i)$
- ▶ This promotes solutions where t_i is large
- ▶ Make sense: we want t_i being large
 - ▶ because $t_i > 0$ means the i th label is correctly predicted



Understanding the Loss

Definition

The function $\ell(t) := \mathbf{1}_{t < 0}$ is called **0-1 loss**. In this context, **1** denotes the indicator function.



Observations:

- ▶ The 0-1 loss is one when the label is misclassified ($t < 0$) and zero otherwise
- ▶ Thus the cumulative 0-1 loss $\sum_{i=1}^n \ell(y_i (\langle \mathbf{w}, \phi(\mathbf{x}_i) \rangle + b))$ measures the number of training errors
- ▶ Makes sense to get this error as small as possible
- ▶ Turns out to be an NP hard problem! :(

Idea in SVM, LR, and ANN: replace the difficult 0-1 loss by a convex approximation—the hinge or logistic loss